

Semaphores Mutexes



M³RTS

by Prof. Dr. rer. nat.
Karsten Weronek

M³Real-Time-Systems

SS 2017

Prof. Dr. rer. nat. Karsten Weronek
Faculty 2
Computer Science and Engineering

Semaphores
Mutexes

For the schedulability tests we had necessary and sufficient conditions.

Question: Is there a feasible schedule?

What are the implications if:

- A necessary condition is violated
- A necessary condition is fulfilled
- A sufficient condition is violated
- A sufficient condition is fulfilled

Possible answers: ①:Yes, ②:No, ③:Perhaps

When you have to solve a Real-Time problem
you probably need a couple of threads or processes that work in parallel.

You have to distinguish the following relationships
between parallel tasks:

- disjunct tasks; they run completely independent
- non-disjunct tasks; they use common data or resources
 - competing tasks, that apply for access to the same data,
 - Tasks that are dependent, because one has deliver for the other (producer/consumer-model)

The result of non-disjunct parallel processes without proper synchronization
is often non-predictable and non-reproducible.

The part of a code sequences (Thread/Process or Interrupt-Service Routines) in which a joint access to common resources may occur are called "Critical Section".

When two or more tasks access common data, this may lead to a "Race Condition".

The result of a Race Condition depends on the relative progress of the different tasks.

1. Race conditions can be avoided by allowing access to a critical section for one task only at a time. This is called **mutual exclusion** (gegenseitiger Ausschluss).
2. The mathematical instrument to solve this issue is called a **semaphore**.

- A semaphore is an integer S
- A semaphore has a Maximum N
- A semaphore is initialized by N
- There is a function P :
 - $S = S - 1$ (passeeren)
 - If $S < 0$ then queue
- There is a function V :
 - $S = S + 1$ (vrijgeven)
 - If $S \leq 0$ then take one out of the queue (highest Prio)

A Semaphore with $N=1$ is called **Binary Semaphore** or a **Mutex**.

It tells you only free or not free and

The size of the queue respectively.

The little room.

- P-Operation:

```
S = S - 1;  
if (s < 0) {  
    sleep_until_semaphore_is_free()  
}
```

- V-Operation:

```
S = S + 1;  
if (s ≤ 0) {  
    wake_up_sleeping_process_with_highest_priority();  
}
```

- P- and V-Operations are critical sections as well and are non-interruptible. Therefore semaphore operations are systemcalls.

```
int pthread_mutex_init (...)  
int pthread_mutex_destroy (...)  
int pthread_mutex_lock (...)  
int pthread_mutex_trylock (...)  
int pthread_mutex_unlock (...)
```

```
#include <pthread.h>                                // compile like: gcc -pthread myProg.c
#define NTHREADS 10
void *thread_function(void *);
pthread_mutex_t mutex1 = PTHREAD_MUTEX_INITIALIZER;
int counter = 0;

main() {
    pthread_t thread_id[NTHREADS];
    int i, j;

    for(i=0; i < NTHREADS; i++){
        pthread_create( &thread_id[i], NULL, thread_function, NULL );
    }
    for(j=0; j < NTHREADS; j++) {
        pthread_join( thread_id[j], NULL);
    }
    printf("Final counter value: %d\n", counter);
}

void *thread_function(void *dummyPtr) {
    printf("Thread number %ld\n", pthread_self());
    pthread_mutex_lock( &mutex1 );
    counter++;
    pthread_mutex_unlock( &mutex1 );
}
```

Let two tasks A and B

Having to Mutexes S1 and S2 the request sequences are:

Task A: P(S1); P(S2); V(S2); V(S1)

Task B: P(S2); P(S1); V(S1); V(S2)

Let two tasks A and B

Having to Mutexes S1 and S2 the request sequences are:

Task A: P(S1); P(S2); V(S2); V(S1)

Task B: P(S2); P(S1); V(S1); V(S2)

Task A	Task B
P(S1)	P(S2)
P(S2)	P(S1)
V(S2)	V(S1)
V(S1)	V(S2)

Deadlock occurs:
S1 is blocked by Task A
S2 is blocked by Task B

There are four sufficient conditions (in combination) for deadlocks:

1. Resources are exclusive and are not accessible by multiple task at a time
2. A task uses at least two resources (e.g. CPU and microphone)
3. The resource can not be withdrawn easily from a using task by another task
4. Meanwhile the usage of resources by different task there exists a cyclic chain

The developer has to avoid only one of the four conditions to avoid deadlocks.

This is not easy!

1. Resources are exclusive and are not accessible by multiple task at a time

This is hard to change. For Data/Memory access it is impossible.
For most of devices or blocks on the board it is also not possible.

2. A task uses at least two resources (e.g. CPU and microphone)

You may isolate pure calculating task in a core.*

However since most of the task have frequent I/O or use Co-Processors or other blocks this is not an option to change.

Another alternative would be if a task always waits until all resources are free. This will take a lot of time....

NB: you can try this in the exercises by calling:

using boot option "Realtime Kernel - Kernel isolation"
isolates CPU1 and CPU2

You may also switch of some cores by using:

```
sudo /opt/m3rts/setCPU1    ...  
sudo /opt/m3rts/setCPU4
```

3. The resource can not be withdrawn easily from a using task by another task

This is dependent of the device.

Most of devices have a data loss when you withdraw it from a task.

4. Meanwhile the usage of resources by different task
there exists a cyclic chain

This can be avoided

by defining a linear order for the access to the required resources. → Piping.

This may be not useful or it may take time.

When the resource is data,

it depends if the task wants to read or to write the data.

because reading data by multiple tasks at the same time is not an issue

but when a task tries to write a data this is an issue for other tasks.

A solution for this issue can be realized by a semaphore

that allows parallel read-request but allows only a single process to write.

See the following cases:

1. The critical section is free
 - Allow access
2. There is on or more tasks, reading the critical section and there is no task to write:
 - A task to read gets access
 - A task to write will be blocked
3. The critical section is accessed by at least one reading task and a writing tasks wants to access
 - Access will be blocked
4. The critical section is accessed by a writing task
 - All other requesting tasks will be blocked.

The preemption model is responsible for the safety of critical sections.

It divides the operating system in different layers.

There are different preemption models and they are implemented differently
(e.g. does a Job preempt another when they have the same priority)

The System Architect needs to know it
because it influences the real-time behavior.

The Application Developer needs to know it
to identify the critical parts of the code and
to be able to assure that they will be secured.

The preemption model is divided into four layers:

- The Application-Layer (Userland)
- The Kernel-Layer
- The Soft-IRQ-Layer
- The Interrupt-Layer

Some Details:

On the application layer user programs runs with certain priorities.

Semaphores are used for critical sections.

When an application starts a thread on the Kernel-Layer it inherits the priority.

User threads compete with Kernel-Threads. Kernel Threads are not interruptible.

Interrupt Service Routines have the highest (all the same) priority (HW-Priority).

On single core machines there is no need to protect data.

After having finished the HW-ISR the SW-IRQ are executed.

Since Interrupts are not allowed for sleep mode, use of semaphores doesn't make sense

The situation where a high priority task requiring a resource is blocked due to the lock of this resource by a low priority task.

Then the high priority task is blocked until the low priority task is completed and has released the resource.

This can take a long time, especially when mid-prioritized tasks delay the completion of the low prioritized one.